

Principles of Data Science – 5530

Assignment-2

Name: Nikhil Jagadeesh Sriram

Student ID: 16352573

(A) In this context, I am addressing the identification and handling of missing values across all columns. For categorical columns, the missing values are being imputed with the mode. This approach is chosen to ensure the preservation of the existing data within that categorical field, thereby minimizing the potential for bias in the imputed values. As for numerical columns, missing values are being imputed with the mean. This method is adopted to maintain the distribution and central tendency of the available data in the numerical column, aligning the imputed values with the overall dataset distribution and mitigating the risk of introducing noticeable bias.

```
missing_values = data_given.isnull().sum()
print(missing_values) # Checking for missing values

miss_col = []
for col, value in missing_values.items():
    if value > 0:
        miss_col.append(col)

print("Columns with missing values:", miss_col)

Unnamed: 0      0
Name            0
Location        0
Year            0
Kilometers_Driven 0
Fuel_Type       0
Transmission    0
Owner_Type      0
Mileage         2
Engine          36
Power           36
Seats           38
New_Price       5032
Price           0
dtype: int64
Columns with missing values: ['Mileage', 'Engine', 'Power', 'Seats', 'New_Price']

[ ] for col in miss_col:
    if data_given[col].dtype == 'object':
        data_given[col].fillna(data_given[col].mode()[0], inplace=True)
    else:
        data_given[col].fillna(data_given[col].mean(), inplace=True)
```

b) Remove the units from some of the attributes and only keep the numerical values (for example remove kmpl from "Mileage", CC from "Engine", bhp from "Power", and lakh from "New_price").

```
column_data_types = car_data.dtypes

# Print the data types of columns
print(column_data_types)

Unnamed: 0      int64
Name           object
Location       object
Year           int64
Kilometers_Driven int64
Fuel_Type      object
Transmission   object
Owner_Type     object
Mileage        object
Engine         object
Power          object
Seats          float64
New_Price      object
Price          float64
dtype: object

[ ] data_given

   Unnamed: 0  Name  Location  Year  Kilometers_Driven  Fuel_Type  Transmission  Owner_Type  Mileage  Engine  Power  Seats  New_Price  Price
0           1  Hyundai Creta 1.6 CRDi SX Option  Pune   2015           41000        Diesel        Manual        First   19.67 kmpl  1582 CC  126.2 bhp   5.0   4.78 Lakh  12.50
```

(B) Strip the units from select attributes, retaining only the numerical values. For instance, eliminate "kmpl" from the "Mileage" column, "CC" from the "Engine" column, "bhp" from the "Power" column, and "lakh" from the "New_Price" column.

5842	6014	Maruti Swift VDI	Delhi	2014	27365	Diesel	Manual	First	28.4 kmpl	1248 CC	74 bhp	5.0	7.88 Lakh	4.75
5843	6015	Hyundai Xcent 1.1 CRDI S	Jaipur	2015	100000	Diesel	Manual	First	24.4 kmpl	1120 CC	71 bhp	5.0	4.78 Lakh	4.00
5844	6016	Mahindra Xylo D4 BSIV	Jaipur	2012	55000	Diesel	Manual	Second	14.0 kmpl	2498 CC	112 bhp	8.0	4.78 Lakh	2.90
5845	6017	Maruti Wagon R VXI	Kolkata	2013	46000	Petrol	Manual	First	18.9 kmpl	998 CC	67.1 bhp	5.0	4.78 Lakh	2.65
5846	6018	Chevrolet Beat Diesel	Hyderabad	2011	47000	Diesel	Manual	First	25.44 kmpl	936 CC	57.6 bhp	5.0	4.78 Lakh	2.50

5847 rows x 14 columns

```
data_given['Mileage'] = data_given['Mileage'].str.extract('(\\d+\\.\\d+)').astype(float)
data_given['Engine'] = data_given['Engine'].str.replace(' CC', '').astype(int)
data_given['Power'] = data_given['Power'].str.extract('(\\d+\\.\\d+)').astype(float)
data_given['New_Price'] = data_given['New_Price'].str.extract('(\\d+\\.\\d+)').astype(float)
```

+ Code + Text

data_given

Unnamed: 0		Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type	Mileage	Engine	Power	Seats	New_Price	Price
0	1	Hyundai Creta 1.6 CRDI SX Option	Pune	2015	41000	Diesel	Manual	First	19.67	1582	126.20	5.0	4.78	12.50
1	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	First	NaN	1199	88.70	5.0	8.61	4.50
2	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	First	20.77	1248	88.76	7.0	4.78	6.00
3	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Second	15.20	1968	140.80	5.0	4.78	17.74
4	6	Nissan Micra Diesel XV	Jaipur	2013	86999	Diesel	Manual	First	23.08	1461	63.10	5.0	4.78	3.50

(C) Convert the categorical variables, namely "Fuel_Type" and "Transmission," into numerical one-hot encoded values.

One-hot encoding is a technique used to convert categorical variables into a binary matrix where each category becomes a separate column, and a binary value (0 or 1) indicates the presence or absence of that category for each observation. In the context of your dataset:

For example, if you have a "Fuel_Type" column with categories 'Petrol' and 'Diesel', after one-hot encoding, you will have two new columns: 'Fuel_Type_Petrol' and 'Fuel_Type_Diesel'. For each row, the column corresponding to the actual fuel type will have a value of 1, while the other column will have a value of 0.

C) Change the categorical variables ("Fuel_Type" and "Transmission") into numerical one hot encoded value.

```
[ ] label_encoder = LabelEncoder()
data_given['Fuel_Type_Label'] = label_encoder.fit_transform(data_given['Fuel_Type'])
data_given['Transmission_Label'] = label_encoder.fit_transform(data_given['Transmission'])

onehot_encoder = OneHotEncoder(sparse=False)

fuel_type_encoded = onehot_encoder.fit_transform(data_given[['Fuel_Type_Label']])
transmission_encoded = onehot_encoder.fit_transform(data_given[['Transmission_Label']])

fuel_type_encoded_df = pd.DataFrame(fuel_type_encoded, columns=['Fuel_Type_' + str(i) for i in range(fuel_type_encoded.shape[1])])
transmission_encoded_df = pd.DataFrame(transmission_encoded, columns=['Transmission_' + str(i) for i in range(transmission_encoded.shape[1])])

data_given = pd.concat([data_given, fuel_type_encoded_df, transmission_encoded_df], axis=1)

data_given.drop(['Fuel_Type', 'Transmission', 'Fuel_Type_Label', 'Transmission_Label'], axis=1, inplace=True)
```

(D) Introduce a new feature by calculating the current age of each car and add this column to the dataset. One way to achieve this is by subtracting the "Year" value from the current year.

d) Create one more feature and add this column to the dataset (you can use mutate function in R for this). For example, you can calculate the current age of the car by subtracting "Year" value from the current year.

current_year = datetime.now().year
data_given["Current_Age"] = current_year - data_given["Year"]
data_given

Unnamed: 0

	Name	Location	Year	Kilometers_Driven	Owner_Type	Mileage	Engine	Power	Seats	New_Price	Price	Fuel_Type_0	Fuel_Type_1	Fuel_Type_2	Transmission_0	Transm
0	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	First	19.67	1582	126.20	5.0	4.78	12.50	1.0	0.0	0.0	0.0
1	2	Honda Jazz V	Chennai	2011	46000	First	NaN	1199	88.70	5.0	8.61	4.50	0.0	0.0	1.0	0.0
2	3	Maruti Ertiga VDI	Chennai	2012	87000	First	20.77	1248	88.76	7.0	4.78	6.00	1.0	0.0	0.0	0.0
3	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Second	15.20	1968	140.80	5.0	4.78	17.74	1.0	0.0	0.0	1.0
4	6	Nissan Micra	Jaipur	2013	86999	First	23.08	1461	63.10	5.0	4.78	3.50	1.0	0.0	0.0	0.0

(E) Perform select, filter, rename, mutate, arrange and summarize with group by operations on this dataset.

In data manipulation, the **select** operation involves choosing specific columns, **filter** is used to subset rows based on conditions, **rename** is employed to change column names, **mutate** creates new columns, **arrange** sort rows based on specified criteria, and **summarize** with **group by** aggregates data by groups, calculating summary statistics for each group.

e) Perform select, filter, rename, mutate, arrange and summarize with group by operations (or their equivalent operations in python) on this dataset.

[] selected_data = data_given[['Name', 'Year', 'Kilometers_Driven']]#select
print(selected_data)

	Name	Year	Kilometers_Driven
0	Hyundai Creta 1.6 CRDi SX Option	2015	41000
1	Honda Jazz V	2011	46000
2	Maruti Ertiga VDI	2012	87000
3	Audi A4 New 2.0 TDI Multitronic	2013	40670
4	Nissan Micra Diesel XV	2013	86999
...
5842	Maruti Swift VDI	2014	27365
5843	Hyundai Xcent 1.1 CRDi S	2015	100000
5844	Mahindra Xylo D4 BSIV	2012	55000
5845	Maruti Wagon R VXI	2013	46000
5846	Chevrolet Beat Diesel	2011	47000

[5847 rows x 3 columns]

filtered_data = data_given[data_given['Location'] == 'Jaipur']#filtered Data
print(filtered_data)

Unnamed: 0

	Name	Location	Year	Kilometers_Driven	\
4	6	Nissan Micra Diesel XV	Jaipur	2013	86999
10	12	Maruti Swift VDI BSIV	Jaipur	2015	64424
15	17	Maruti Swift DDiS VDI	Jaipur	2017	25000

[403 rows x 18 columns]

```
data_given = data_given.rename(columns={'New_Price': 'NewPriceCAR'})#rename
data_given
```

Unnamed: 0 Name Location Year Kilometers_Driven Owner_Type Mileage Engine Power Seats N

		Hyundai										
--	--	---------	--	--	--	--	--	--	--	--	--	--

Unnamed: 0

5847 rows x 18 columns

```
[ ] data_given['Increseprice'] = data_given['NewPriceCAR']- data_given['Price']
data_given
```

Unnamed: 0 Name Location Year Kilometers_Driven Owner_Type Mileage Engine Power Seats NewPriceCAR Price Fuel_Ty

0	1	Hyundai Creta 1.6 CRDI SX Option	Pune	2015		41000	First	19.67	1582	126.20	5.0	4.78	12.50
		..											

5847 rows x 19 columns

```
sorted_car_data = car_data.sort_values(by='Price', ascending=True)
sorted_car_data#arrange
```

Unnamed: 0 Name Location Year Kilometers_Driven Owner_Type Mileage Engine Power Seats Price

1660	1713	Tata Nano LX	Pune	2011		65000	Second	26.00	624	
------	------	-----------------	------	------	--	-------	--------	-------	-----	--

5847 rows x 19 columns

```
average_price_by_fuel_type = data_given.groupby('Kilometers_Driven')['Mileage'].mean().#summarize
average_price_by_fuel_type
```

Kilometers_Driven

171	24.700000
600	21.500000
1000	18.493333
1001	22.415000
1011	20.370000
...	
480000	17.180000
620000	20.360000
720000	20.540000
775000	19.300000
6500000	15.970000

Name: Mileage, Length: 3019, dtype: float64

```
[ ] data_given.to_csv('clean_data.csv', index=False)
```